Yue Gao

📱 +1 (608) 733-8789   ✉ gy@cs.wisc.edu   🏠 pages.cs.wisc.edu/~gy   💼 ygao234   🐙 Lodour   📍 Madison, WI

RESEARCH INTERESTS

AI Security	ML System Security, Adversarial Robustness, Attacks and Defenses, Security and Privacy.
Machine Learning	Vision Recognition, Natural Language Processing, Multi-Modality, Diffusion Models, LLMs.
Cybersecurity	Multi-component ML Systems, Web-based Applications, Cryptography, Linux Kernel Memory.

EDUCATION

University of Wisconsin–Madison   Madison, WI
Ph.D. Candidate in Computer Science (advised by Prof. Kassem Fawaz)   *Sep 2018 – Jan 2024 (expected)*
- Thesis: *Characterizing the Limitations of Defenses in Adversarial Machine Learning*
- Selected Courses: *Introduction to Information Security, Applied Cryptography, Advanced Operating Systems.*

Shanghai University   Shanghai, China
B.S. in Computer Science and Technology (GPA 3.99/4.00, Ranked 1/292)   *Sep 2014 – Jul 2018*
- Thesis: *A Deep Neural Network based Image Compression Method*
- Selected Courses: *Operating Systems, Computer Network, Assembly Language, Software Engineering.*

WORK EXPERIENCE

Research Assistant @ University of Wisconsin–Madison   Madison, WI
Advised by Prof. Kassem Fawaz   *Nov 2018 – Present*
- Researched the security of multi-component ML Systems deployed in real-world environments.
- Systematized the security analysis of ML-based and web-based systems in black-box settings.

ML System Security Research Intern @ Microsoft Research   Redmond, WA
Mentored by Dr. Jay Stokes and Dr. Emre Kiciman   *Jun 2021 – Sep 2021*
- Proposed a research project on defenses against imperceptible textual backdoor attacks on language models.
- Discovered blind spots in state-of-the-art attacks and defenses, and published stronger defenses at MILCOM.
- Successfully reduced the attack success rate from 100% to 12%, even at a challenging poisoning rate of 10%.

ML Research and Development Intern @ TuCodec (Startup)   Shanghai, China
Mentored by Dr. Chunlei Cai   *Jan 2018 – Jul 2018*
- Secured 1st place as a primary contributor in the CVPR 2018 Challenge on Learned Image Compression.
- Improved the average runtime of DNN-based image compression algorithms from 1 min to 4 secs per 4K-res image.
- Independently developed ML-based desktop apps on Ubuntu, MacOS, and Windows using TensorFlow and C++.
- Independently developed ML-based image compression systems using TensorFlow, Python, Docker, and Kubernetes.
- Protected ML systems from security intrusion and model stealing with security measures and anomaly analysis.

PUBLICATIONS

Conference

[1] **On the Limitations of Stochastic Pre-processing Defenses**
Yue Gao, Ilia Shumailov, Kassem Fawaz, and Nicolas Papernot
Proceedings of the 36th Conference on Neural Information Processing Systems (NeurIPS), 2022

[2] **Rethinking Image-Scaling Attacks: The Interplay Between Vulnerabilities in Machine Learning Systems**
Yue Gao, Ilia Shumailov, and Kassem Fawaz
Proceedings of the 39th International Conference on Machine Learning (ICML), 2022
Oral Presentation (Top 2%)

[3] **Experimental Security Analysis of the App Model in Business Collaboration Platforms**
Yunang Chen*, *Yue Gao**, Nick Ceccio, Rahul Chatterjee, Kassem Fawaz, and Earlence Fernandes
31st USENIX Security Symposium (USENIX Security), 2022
Bug Bounty ($1500)

[4] **I Know Your Triggers: Defending Against Textual Backdoor Attacks With Benign Backdoor Augmentation**
Yue Gao, Jack W. Stokes, Manoj Prasad, Andrew Marshall, Kassem Fawaz, and Emre Kiciman
IEEE Military Communications Conference (MILCOM), 2022

Workshop

[1] **Variational Autoencoder for Low Bit-rate Image Compression**
Lei Zhou*, Chunlei Cai*, *Yue Gao*, Sanbao Su, and Junmin Wu
Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR) Workshops, 2018
Winner of the first Challenge on Learned Image Compression

Preprints

[1] **SEA: Shareable and Explainable Attribution for Query-based Black-box Attacks**
Yue Gao, Ilia Shumailov, and Kassem Fawaz
arXiv, 2023

[2] **Human-Producible Adversarial Examples**
David Khachaturov, *Yue Gao*, Ilia Shumailov, Robert Mullins, Ross Anderson, and Kassem Fawaz
arXiv, 2023

[3] **Analyzing Accuracy Loss in Randomized Smoothing Defenses**
Yue Gao, Harrison Rosenberg*, Kassem Fawaz, Somesh Jha, and Justin Hsu
arXiv, 2020

SELECTED PROJECTS

Security of Real-world ML Systems: Threat Modeling, Defending, and Characterization   *Sep 2020 – May 2022*

- Investigated the security of multi-component ML systems exposed to diverse security threats, e.g., dependencies.
- Revealed threats amplified by 9x and broke state-of-the-art defenses by jointly exploiting multiple vulnerabilities.
- Formally proved the non-robustness of randomization-based defenses beyond demonstrating empirical attacks.
- Characterized the attack's progression for forensic purposes and human-explainable intelligence sharing.

Guaranteeing AI Robustness Against Deception (DARPA Engagement)   *Jun 2019 – Feb 2024*

- Led a 9-member cross-university team to 1st and 2nd place in grant competitions hosted by DARPA.
- Performed red teaming and broke over 10 internal defenses proposed by team members prior to submission.
- Proposed CLIP-like and diffusion-based methods to enforce robust features across RGB and Depth modalities.
- Successfully reduced the disappearance rate from 62% to 9% even under the red-team evaluation from MITRE.
- Developed initial code bases and eval pipelines for team members from varying technical backgrounds.
- Contributed plug-and-play modules to the official upstream evaluation team and received acknowledgment.

Security Analysis of Online Business Collaboration Platforms   *Mar 2021 – Dec 2021*

- Automated the analysis of security principle violations for 3000+ third-party apps in Slack and Microsoft Teams.
- Reverse engineered OAuth designs to bypass access control, and received bug bounty for medium severity.
- Demonstrated POC attacks of eavesdropping on private chats, spoofing video calls, and unauthorized code merging.

SELECTED HONORS & AWARDS

Slack Bug Bounty: Medium Severity, $1500	2022
Top 10% Reviewers Award: NeurIPS	2022
CVPR Competition Winner: Challenge on Learned Image Compression	2018
National Scholarship: China	2017
Top 100 Elite Collegiate Award: China Computer Federation	2017
Scholarship for Exceptional Leadership: Shanghai University	2017
City Scholarship: Shanghai	2016
Outstanding Student Award: Shanghai University	2016
Outstanding Volunteer Award: ACM ICPC Asia Regional Contest	2016
Scholarship for Exceptional Innovation: Shanghai University	2016

Scholarship for Exceptional Academic Achievements: Shanghai University	2015 – 2018
Bronze Prize for Programming Contest: ACM ICPC Asia East-Continent Final Contest	2015
Bronze Prize for Programming Contest: ACM ICPC Asia Shanghai Regional Contest	2015

PROFESSIONAL ACTIVITIES

Reviewer: NeurIPS, ICML, and ICLR	2022 – 2024
External Reviewer: IEEE Symposium on Security and Privacy	2021 – 2023
External Reviewer: USENIX Security Symposium	2021 – 2022
External Reviewer: ACM Conference on Computer and Communications Security	2019
Team Leader: Collegiate ICPC Team at Shanghai University	2016 – 2017

TALKS

1. **Forensics and Intelligence Sharing for ML Security** *Oct 2023*
   DARPA GARD PI Meeting, IBM Research
2. **The Vulnerabilities of Preprocessing in Adversarial Machine Learning** *Oct 2023*
   ML Red Team, Google
3. **The Vulnerabilities of Preprocessing in Adversarial Machine Learning** *Apr 2023*
   TrustML Young Scientist Seminar, RIKEN AIP
4. **On the Limitations of Stochastic Pre-processing Defenses** *Oct 2022*
   DARPA GARD PI Meeting, University of Southern California (virtual)
5. **The Interplay Between Vulnerabilities in Machine Learning Systems** *Sep 2022*
   DARPA GARD PI Meeting, University of Michigan
6. **Experimental Security Analysis of the App Model in Business Collaboration Platforms** *Aug 2022*
   USENIX Security 2022
7. **The Interplay Between Vulnerabilities in Machine Learning Systems** *Jun 2022*
   ICML 2022

TEACHING AND MENTORING

Project Mentor: DARPA GARD Project, University of Wisconsin–Madison	Fall 2023
Teaching Assistant: CS 368 (C++ for Java Programmers), University of Wisconsin–Madison	Fall 2018
Guest Lecturer: Advanced Algorithms & Data Structures, Shanghai University	2015 – 2017
Problem Designer: Undergraduate Programming Contests, Shanghai University	2015 – 2017
Student Mentor: Undergraduate Computer Science Coursework, Shanghai University	2015 – 2017

TECHNICAL SKILLS (ORDERED BY EXPERTISE)

Python	ML and Security Research (2018 – present), Backend Development (2016 – 2018).
PyTorch	ML Research (2019 – present), Distributed Training (2020 – present).
Docker	Research (2018 – present), Computing Cluster (2017 – present), Model Deployment (2017 – 2018).
C / C++	Linux Kernel (2019), Encryption (2019), Software Development (2017 - 2018), ICPC (2014 – 2018).
Security	CTF (2015 – 2017, with IDA Pro, OllyDbg, Burp Suite, and nmap).
TensorFlow	ML Research (2017 – 2020), Service Deployment (2018).
Java EE	Backend Development (2016).

ARTICLES AND MEDIA COVERAGE

CleverHans. Can stochastic pre-processing defenses protect your models?	2022
USENIX login. Experimental Security Analysis of the App Model in Business Collaboration Platforms	2022
Wired. Slack's and Teams' Lax App Security Raises Alarms	2022